

CEG 4350/6350 Final Exam. Summer 2020
Open Book, notes, and handouts (9:50 – 11:55 am)
Alex Reigle

1. (15 points)

Suppose there are two types of resources – resource type R and resource type S . Resource type R has m identical instances, and resource type S has n identical instances. A number of processes can share these instances as follows:

1. Each instance (of resource types R and S) can be used by only one process at a time.
2. Each process needs to acquire and use any one instance of R and any one instance of S at the same time. Thus, if there is no instance of resource type R or S (or both) available for a requesting process, the process has to wait until both instances become available.

Design a monitor implementing `Acquire()` and `Release()` functions that behave as follows (in atomic manner); and declare and initialize as many variables as needed.

1.

Monitor ResourceAllocate{

//Monitors by definition satisfy mutual exclusion

Semaphore semS;

Semaphore semR;

semS.S = n; //Initialize to max available number of resources/instances

semR.R = m;

Acquire () {

if (semR.R == 0) {

block(semR); //block this process in a queue

}

if(semS.S == 0){

block(semS); //block this process in a queue

}

/*Wait until there is a resource available and
then use that resource*/

semR.R--;

semS.S--;

// return the ids of two available resource instances (note: this comment doesn't need to be
// implemented)

}

Release () {

// mark two released resource instances as available (note: this comment doesn't need to be
// implemented)

if(semR.R < m){

wakeup(semR); //wake up a blocked process (if there is any);

semR.R++;}

if(semS.S < n){

wakeup(semS); //wake up a blocked process (if there is any);

semS.S++;}

/* If the semaphore values are equal to their max (m or n) then there are no resources to
release and the program should exit with no action*


}

}

2. (15 points)

Consider two periodic processes, P1 and P2, whose period and processing time are $(p_1=50, t_1=25)$ and $(p_2=75, t_2=30)$, respectively. Note: As both P1 and P2 are periodic processes, each process will start only once within each of its period.

- (1) Can these two processes be scheduled using the rate-monotonic scheduling (RMS)? Justify your answer using a Gantt chart (i.e., a timing diagram).




Yes, RMS can be used to schedule these processes. However, the RMS is not guaranteed to work based on:

$$\sum_{i=1}^n \frac{t_i}{p_i} \leq n(2^{1/n} - 1)$$

Where the sum of $t_i/p_i = 0.9$ which is greater than 0.83 and therefor does not satisfy the guarantee condition.

- (2) Can these two processes be scheduled using the earliest-deadline-first (EDF) scheduling? Justify your answer using a Gantt chart (i.e., a timing diagram).



Yes, EDF can be used to schedule these processes.

3. (15 points)

Suppose that 4 processes, denoted by P1, P2, P3 and P4, are sharing the memory that has total 16 frames (numbered from 0 to 15), and currently some of their virtual pages are loaded into the frames as follows:

- Process P1's virtual pages 100, 101, 102, and 103 are loaded at frames 4, 5, 6, and 7, respectively.
- Process P2's virtual pages 200, 201, 202, and 203 are loaded at frames 0, 1, 2, and 3, respectively.
 - Process P3's virtual pages 300, 301, 302, and 303 are loaded at frames 12, 13, 14, and 15, respectively.
- Process P4's virtual pages 400, 401, 402, and 403 are loaded at frames 8, 9, 10, and 11, respectively.

We also assume that an **inverted page table (IPT)** is used for address translation.

- (a) Show the structure (i.e., entry numbers and column names) and the current content (i.e., current values in each entry) of the **inverted page table**.

INVERTED PAGE TABLE			
Frame #	Page #	RWX	Dirty Bit
0	200		
1	201		
2	202		
3	203		
4	100		
5	101		
6	102		
7	103		
8	400		
9	401		
10	402		
11	403		
12	300		
13	301		
14	302		
15	303		

- (b) Show the structure and current content of the **hash table** created additionally, in order to speed up the searching of the above inverted page table. We assume the hash function used is: virtual_page_number *modulo* 10 (i.e., virtual_page_number % 10).

HASH TABLE	
Frame modulo(10)	Frames
0	0 --> 4 --> 8 --> 12
1	1 --> 5 --> 9 --> 13
2	2 --> 6 --> 10 --> 14
3	3 --> 7 --> 11 --> 15
4	
5	
6	
7	
8	
9	